

#CROP YIELD PREDICTION

#Dataset 1

##Reading Dataset

```
In [2]: 1 import pandas as pd
        2 df = pd.read_csv('finalised_dataset.csv',na_values='')
        3 df
```

Out[2]:

	Unnamed: 0	state_names	district_names	crop_year	season_names	crop_names	are
0	125191	Maharashtra	AHMEDNAGAR	1997	Autumn	Maize	1.
1	125192	Maharashtra	AHMEDNAGAR	1997	Kharif	Arhar/Tur	17600.
2	125193	Maharashtra	AHMEDNAGAR	1997	Kharif	Bajra	274100.
3	125194	Maharashtra	AHMEDNAGAR	1997	Kharif	Gram	40800.
4	125195	Maharashtra	AHMEDNAGAR	1997	Kharif	Jowar	900.
...
12623	137814	Maharashtra	YAVATMAL	2014	Rabi	Jowar	4000.
12624	137815	Maharashtra	YAVATMAL	2014	Rabi	Maize	1300.
12625	137816	Maharashtra	YAVATMAL	2014	Rabi	Wheat	29100.
12626	137817	Maharashtra	YAVATMAL	2014	Summer	Groundnut	9400.
12627	137818	Maharashtra	YAVATMAL	2014	Whole Year	Sugarcane	8100.

12628 rows × 17 columns

```
In [3]: 1 df=df.drop('Yield', axis = 1)
```

In [4]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12628 entries, 0 to 12627
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      12628 non-null  int64
1   state_names     12628 non-null  object
2   district_names  12628 non-null  object
3   crop_year       12628 non-null  int64
4   season_names    12628 non-null  object
5   crop_names      12628 non-null  object
6   area            12628 non-null  float64
7   temperature     12628 non-null  float64
8   wind_speed      12628 non-null  float64
9   pressure        12628 non-null  float64
10  humidity        12628 non-null  float64
11  soil_type       12628 non-null  object
12  N               12628 non-null  float64
13  P               12628 non-null  float64
14  K               12628 non-null  float64
15  production      12496 non-null  float64
dtypes: float64(9), int64(2), object(5)
memory usage: 1.5+ MB
```

In [5]: 1 df.columns
2

Out[5]: Index(['Unnamed: 0', 'state_names', 'district_names', 'crop_year',
'season_names', 'crop_names', 'area', 'temperature', 'wind_speed',
'pressure', 'humidity', 'soil_type', 'N', 'P', 'K', 'production'],
dtype='object')

##Reducing Data to One State for Ease

```
In [6]: 1 df = df[df['state_names'] == "Maharashtra"]
        2
        3 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 12628 entries, 0 to 12627
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Unnamed: 0      12628 non-null  int64
1   state_names     12628 non-null  object
2   district_names  12628 non-null  object
3   crop_year       12628 non-null  int64
4   season_names    12628 non-null  object
5   crop_names      12628 non-null  object
6   area            12628 non-null  float64
7   temperature     12628 non-null  float64
8   wind_speed      12628 non-null  float64
9   pressure        12628 non-null  float64
10  humidity        12628 non-null  float64
11  soil_type       12628 non-null  object
12  N               12628 non-null  float64
13  P               12628 non-null  float64
14  K               12628 non-null  float64
15  production      12496 non-null  float64
dtypes: float64(9), int64(2), object(5)
memory usage: 1.6+ MB
```

```
In [7]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 12628 entries, 0 to 12627
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Unnamed: 0      12628 non-null  int64
1   state_names     12628 non-null  object
2   district_names  12628 non-null  object
3   crop_year       12628 non-null  int64
4   season_names    12628 non-null  object
5   crop_names      12628 non-null  object
6   area            12628 non-null  float64
7   temperature     12628 non-null  float64
8   wind_speed      12628 non-null  float64
9   pressure        12628 non-null  float64
10  humidity        12628 non-null  float64
11  soil_type       12628 non-null  object
12  N               12628 non-null  float64
13  P               12628 non-null  float64
14  K               12628 non-null  float64
15  production      12496 non-null  float64
dtypes: float64(9), int64(2), object(5)
memory usage: 1.6+ MB
```

```
In [8]: 1 df.isnull().sum()
```

```
Out[8]: Unnamed: 0      0
state_names      0
district_names   0
crop_year        0
season_names     0
crop_names       0
area            0
temperature      0
wind_speed       0
pressure         0
humidity         0
soil_type        0
N               0
P               0
K               0
production      132
dtype: int64
```

```
In [9]: 1 df.head(6)
```

```
Out[9]:
```

	Unnamed: 0	state_names	district_names	crop_year	season_names	crop_names	area	te
0	125191	Maharashtra	AHMEDNAGAR	1997	Autumn	Maize	1.0	
1	125192	Maharashtra	AHMEDNAGAR	1997	Kharif	Arhar/Tur	17600.0	
2	125193	Maharashtra	AHMEDNAGAR	1997	Kharif	Bajra	274100.0	
3	125194	Maharashtra	AHMEDNAGAR	1997	Kharif	Gram	40800.0	
4	125195	Maharashtra	AHMEDNAGAR	1997	Kharif	Jowar	900.0	
5	125196	Maharashtra	AHMEDNAGAR	1997	Kharif	Maize	4400.0	

```
###Making Yield Column
```

In [10]:

1 df

Out[10]:

	Unnamed: 0	state_names	district_names	crop_year	season_names	crop_names	are
0	125191	Maharashtra	AHMEDNAGAR	1997	Autumn	Maize	1.
1	125192	Maharashtra	AHMEDNAGAR	1997	Kharif	Arhar/Tur	17600.
2	125193	Maharashtra	AHMEDNAGAR	1997	Kharif	Bajra	274100.
3	125194	Maharashtra	AHMEDNAGAR	1997	Kharif	Gram	40800.
4	125195	Maharashtra	AHMEDNAGAR	1997	Kharif	Jowar	900.
...
12623	137814	Maharashtra	YAVATMAL	2014	Rabi	Jowar	4000.
12624	137815	Maharashtra	YAVATMAL	2014	Rabi	Maize	1300.
12625	137816	Maharashtra	YAVATMAL	2014	Rabi	Wheat	29100.
12626	137817	Maharashtra	YAVATMAL	2014	Summer	Groundnut	9400.
12627	137818	Maharashtra	YAVATMAL	2014	Whole Year	Sugarcane	8100.

12628 rows × 16 columns


##Corelation Heatmap

In [11]:

```

1 import matplotlib.pyplot as plt
2 import seaborn as sb
3
4 C_mat = df.corr()
5 fig = plt.figure(figsize = (15,15))
6
7 sb.heatmap(C_mat, vmax = .8, square = True,cmap='Blues',annot=True)
8 plt.show()

```



##Taking Data only after 2004 because the data after 2004 is affecting the result alot

```
In [12]: 1 df = df[df['crop_year']>=2004]
          2 df
```

Out[12]:

	Unnamed: 0	state_names	district_names	crop_year	season_names	crop_names	are
212	125403	Maharashtra	AHMEDNAGAR	2004	Kharif	Arhar/Tur	12200.
213	125404	Maharashtra	AHMEDNAGAR	2004	Kharif	Bajra	240500.
214	125405	Maharashtra	AHMEDNAGAR	2004	Kharif	Groundnut	5300.
215	125406	Maharashtra	AHMEDNAGAR	2004	Kharif	Jowar	100.
216	125407	Maharashtra	AHMEDNAGAR	2004	Kharif	Maize	11400.
...
12623	137814	Maharashtra	YAVATMAL	2014	Rabi	Jowar	4000.
12624	137815	Maharashtra	YAVATMAL	2014	Rabi	Maize	1300.
12625	137816	Maharashtra	YAVATMAL	2014	Rabi	Wheat	29100.
12626	137817	Maharashtra	YAVATMAL	2014	Summer	Groundnut	9400.
12627	137818	Maharashtra	YAVATMAL	2014	Whole Year	Sugarcane	8100.

7255 rows × 16 columns

```
In [13]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 7255 entries, 212 to 12627
Data columns (total 16 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      7255 non-null  int64
1   state_names     7255 non-null  object
2   district_names  7255 non-null  object
3   crop_year       7255 non-null  int64
4   season_names    7255 non-null  object
5   crop_names      7255 non-null  object
6   area            7255 non-null  float64
7   temperature     7255 non-null  float64
8   wind_speed      7255 non-null  float64
9   pressure        7255 non-null  float64
10  humidity        7255 non-null  float64
11  soil_type       7255 non-null  object
12  N               7255 non-null  float64
13  P               7255 non-null  float64
14  K               7255 non-null  float64
15  production      7143 non-null  float64
dtypes: float64(9), int64(2), object(5)
memory usage: 963.6+ KB
```

##converting data to numerical form

```
In [14]: 1 df = df.join(pd.get_dummies(df['district_names']))
2 df = df.join(pd.get_dummies(df['season_names']))
3 df = df.join(pd.get_dummies(df['crop_names']))
4 df = df.join(pd.get_dummies(df['state_names']))
5 df = df.join(pd.get_dummies(df['soil_type']))
6 df
```

Out[14]:

	Unnamed: 0	state_names	district_names	crop_year	season_names	crop_names	are
212	125403	Maharashtra	AHMEDNAGAR	2004	Kharif	Arhar/Tur	12200.
213	125404	Maharashtra	AHMEDNAGAR	2004	Kharif	Bajra	240500.
214	125405	Maharashtra	AHMEDNAGAR	2004	Kharif	Groundnut	5300.
215	125406	Maharashtra	AHMEDNAGAR	2004	Kharif	Jowar	100.
216	125407	Maharashtra	AHMEDNAGAR	2004	Kharif	Maize	11400.
...
12623	137814	Maharashtra	YAVATMAL	2014	Rabi	Jowar	4000.
12624	137815	Maharashtra	YAVATMAL	2014	Rabi	Maize	1300.
12625	137816	Maharashtra	YAVATMAL	2014	Rabi	Wheat	29100.
12626	137817	Maharashtra	YAVATMAL	2014	Summer	Groundnut	9400.
12627	137818	Maharashtra	YAVATMAL	2014	Whole Year	Sugarcane	8100.

7255 rows × 88 columns

```
In [15]: 1 df['Yield'] = df['production']/df['area']
2 df
```

Out[15]:

	Unnamed: 0	state_names	district_names	crop_year	season_names	crop_names	are
212	125403	Maharashtra	AHMEDNAGAR	2004	Kharif	Arhar/Tur	12200.
213	125404	Maharashtra	AHMEDNAGAR	2004	Kharif	Bajra	240500.
214	125405	Maharashtra	AHMEDNAGAR	2004	Kharif	Groundnut	5300.
215	125406	Maharashtra	AHMEDNAGAR	2004	Kharif	Jowar	100.
216	125407	Maharashtra	AHMEDNAGAR	2004	Kharif	Maize	11400.
...
12623	137814	Maharashtra	YAVATMAL	2014	Rabi	Jowar	4000.
12624	137815	Maharashtra	YAVATMAL	2014	Rabi	Maize	1300.
12625	137816	Maharashtra	YAVATMAL	2014	Rabi	Wheat	29100.
12626	137817	Maharashtra	YAVATMAL	2014	Summer	Groundnut	9400.
12627	137818	Maharashtra	YAVATMAL	2014	Whole Year	Sugarcane	8100.

7255 rows × 89 columns


```
In [16]: 1 df = df.drop('production', axis=1)
```

##Dropping Unecessary Columns

```
In [17]: 1 df=df.drop('district_names', axis=1)
2 df = df.drop('season_names',axis=1)
3 df = df.drop('crop_names',axis=1)
4
5
6
```

```
In [18]: 1 df = df.drop('state_names', axis=1)
2 df = df.drop('soil_type', axis=1)
3 df
```

Out[18]:

	Unnamed: 0	crop_year	area	temperature	wind_speed	pressure	humidity	
212	125403	2004	12200.0	20.768143	2.002031	1013.280471	20.427922	10.500
213	125404	2004	240500.0	20.722713	2.105239	1015.061641	20.468584	39.720
214	125405	2004	5300.0	21.419190	2.046843	1015.770055	21.836158	8.000
215	125406	2004	100.0	20.425919	2.024060	1013.971163	21.028403	5.820
216	125407	2004	11400.0	20.823344	1.989898	1015.453191	20.340815	0.000
...
12623	137814	2014	4000.0	21.635879	2.000060	1014.302213	20.060662	7.840
12624	137815	2014	1300.0	21.709611	2.053609	1015.803912	21.263478	7.500
12625	137816	2014	29100.0	21.851730	2.027476	1014.031903	20.059945	0.000
12626	137817	2014	9400.0	21.569380	2.004421	1013.989125	21.835158	2.890
12627	137818	2014	8100.0	21.666723	2.008003	1015.081619	21.754799	3.150

7255 rows × 83 columns

##Preprocessing

```
In [19]: 1 from sklearn import preprocessing
```

In [20]:

```

1
2 # Create x, where x the 'scores' column's values as floats
3 x = df[['area']].values.astype(float)
4 x
5 # Create a minimum and maximum processor object
6 min_max_scaler = preprocessing.MinMaxScaler()
7
8 # Create an object to transform the data to fit minmax processor
9 x_scaled = min_max_scaler.fit_transform(x)
10
11 # Run the normalizer on the dataframe
12 #df_normalized = pd.DataFrame(x_scaled)
13 x_scaled
14
15 df['area'] = x_scaled
16 df

```

Out[20]:

	Unnamed: 0	crop_year	area	temperature	wind_speed	pressure	humidity	N
212	125403	2004	0.017150	20.768143	2.002031	1013.280471	20.427922	10.500
213	125404	2004	0.338112	20.722713	2.105239	1015.061641	20.468584	39.720
214	125405	2004	0.007450	21.419190	2.046843	1015.770055	21.836158	8.008
215	125406	2004	0.000139	20.425919	2.024060	1013.971163	21.028403	5.824
216	125407	2004	0.016026	20.823344	1.989898	1015.453191	20.340815	0.000
...
12623	137814	2014	0.005622	21.635879	2.000060	1014.302213	20.060662	7.840
12624	137815	2014	0.001826	21.709611	2.053609	1015.803912	21.263478	7.504
12625	137816	2014	0.040910	21.851730	2.027476	1014.031903	20.059945	0.000
12626	137817	2014	0.013214	21.569380	2.004421	1013.989125	21.835158	2.890
12627	137818	2014	0.011386	21.666723	2.008003	1015.081619	21.754799	3.150

7255 rows × 83 columns

In [21]:

```
1 df.head()
```

Out[21]:

	Unnamed: 0	crop_year	area	temperature	wind_speed	pressure	humidity	N
212	125403	2004	0.017150	20.768143	2.002031	1013.280471	20.427922	10.500
213	125404	2004	0.338112	20.722713	2.105239	1015.061641	20.468584	39.720
214	125405	2004	0.007450	21.419190	2.046843	1015.770055	21.836158	8.008
215	125406	2004	0.000139	20.425919	2.024060	1013.971163	21.028403	5.824
216	125407	2004	0.016026	20.823344	1.989898	1015.453191	20.340815	0.000

5 rows × 83 columns

Filling Empty Values With Mean

```
In [22]: 1 df = df.fillna(df.mean())
```

##Train and Test Split

```
In [23]: 1 from sklearn.model_selection import train_test_split
```

```
In [24]: 1 a=df
```

```
In [25]: 1 b = df['Yield']  
2 #a = df.drop('Yield', axis = 1)  
3  
4
```

```
In [26]: 1 c = df.drop('Unnamed: 0', axis = 1)  
2
```

```
In [:] 1
```

```
In [27]: 1 a=c.drop('Yield', axis = 1)
```

```
In [28]: 1 len(a.columns)
```

Out[28]: 81

```
In [29]: 1 a.columns
```

Out[29]: Index(['crop_year', 'area', 'temperature', 'wind_speed', 'pressure',
'humidity', 'N', 'P', 'K', 'AHMEDNAGAR', 'AKOLA', 'AMRAVATI',
'AURANGABAD', 'BEED', 'BHANDARA', 'BULDHANA', 'CHANDRAPUR', 'DHULE',
'GADCHIROLI', 'GONDIA', 'HINGOLI', 'JALGAON', 'JALNA', 'KOLHAPUR',
'LATUR', 'NAGPUR', 'NANDED', 'NANDURBAR', 'NASHIK', 'OSMANABAD',
'PALGHAR', 'PARBHANI', 'PUNE', 'RAIGAD', 'RATNAGIRI', 'SANGLI',
'SATARA', 'SINDHUDURG', 'SOLAPUR', 'THANE', 'WARDHA', 'WASHIM',
'YAVATMAL', 'Kharif', 'Rabi', 'Summer', 'Whole Year',
'Arhar/Tur', 'Bajra', 'Castor seed', 'Cotton(lint)', 'Gram',
'Groundnut', 'Jowar', 'Linseed', 'Maize', 'Moong(Green Gram)',
'Niger seed', 'Other Rabi pulses', 'Other Cereals & Millets',
'Other Kharif pulses', 'Ragi', 'Rapeseed & Mustard', 'Rice', 'Safflower',
'Sesamum', 'Soyabean', 'Sugarcane', 'Sunflower', 'Tobacco', 'Urad',
'Wheat', 'other oilseeds', 'Maharashtra', 'chalky', 'clay', 'loamy',
'peaty', 'sandy', 'silt', 'silty'],
dtype='object')


```
In [35]: 1 a_train, a_test, b_train, b_test = train_test_split(a, b, test_size = 0.3, random_state = 42)
2
3 print(a_train)
4 print(a_test)
5 print(b_train)
6 print(b_test)
```

	crop_year	area	temperature	wind_speed	pressure	humidity \
8196	2008	0.058343	21.287581	2.066838	1014.904551	20.709584
1674	2014	0.058764	21.080175	2.103093	1014.113752	20.359939
7662	2005	0.013354	20.922660	2.022428	1013.501715	20.642090
2732	2009	0.002389	20.331978	2.019252	1014.653555	21.467039
12588	2012	0.083789	20.773241	2.099441	1013.971550	21.542199
...
6581	2014	0.000702	20.090749	1.983827	1013.157407	20.411220
9120	2006	0.007731	21.544164	2.089016	1013.109528	21.377246
9155	2008	0.003935	20.960905	2.006062	1013.565065	21.525920
9319	2014	0.000420	21.965921	2.061763	1014.285744	21.873364
1582	2010	0.015604	20.139821	2.106225	1013.705632	20.377305

	N	P	K	AHMEDNAGAR ...	Wheat	other oilseeds \
8196	38.250	38.250	38.25	0 ...	0	0
1674	0.000	0.000	0.00	0 ...	0	0
7662	1.600	2.000	0.00	0 ...	0	0
2732	0.500	0.000	0.00	0 ...	0	0
12588	0.000	14.128	0.00	0 ...	0	0
...
6581	0.224	0.280	0.00	0 ...	0	0
9120	0.000	26.832	0.00	0 ...	0	0
9155	7.588	7.588	0.00	0 ...	0	0
9319	2.921	2.921	0.00	0 ...	0	0
1582	7.476	7.476	0.00	0 ...	0	0

	Maharashtra	chalky	clay	loamy	peaty	sandy	silt	silty
8196	1	0	1	0	0	0	0	0
1674	1	0	0	0	0	0	1	0
7662	1	0	0	0	0	0	1	0
2732	1	0	1	0	0	0	0	0
12588	1	0	0	0	0	0	0	1
...
6581	1	0	0	0	0	0	1	0
9120	1	0	0	0	1	0	0	0
9155	1	0	0	0	0	0	1	0
9319	1	0	1	0	0	0	0	0
1582	1	0	0	1	0	0	0	0

[5078 rows x 81 columns]

	crop_year	area	temperature	wind_speed	pressure	humidity \
5229	2010	0.000280	21.626783	2.093701	1015.509013	20.251363
3067	2009	0.044143	20.390844	2.036080	1015.153060	20.485877
6937	2010	0.000420	20.764626	1.999956	1014.064322	21.919638
11148	2005	0.000702	21.920160	2.047580	1014.961353	20.462800
11818	2005	0.000842	22.004754	1.993841	1015.757024	20.567282
...
5269	2012	0.406297	20.091516	2.095636	1014.099868	21.299885
3845	2012	0.007309	20.721014	1.991891	1013.609739	20.340941

2394	2011	0.000139	21.757283	1.973096	1013.299136	21.814687
12214	2011	0.086601	20.883769	2.045719	1015.076670	20.170742
9475	2007	0.000014	21.392058	2.033645	1014.615628	21.072524

	N	P	K	AHMEDNAGAR	...	Wheat	other oilseeds \
5229	1.500	0.000	0.00	0	...	0	0
3067	5.980	5.980	0.00	0	...	0	0
6937	8.740	8.740	0.00	0	...	0	0
11148	3.496	3.496	0.00	0	...	0	0
11818	0.000	0.000	0.00	0	...	0	0
...
5269	38.250	38.250	38.25	0	...	0	0
3845	2.775	0.000	0.00	0	...	0	0
2394	0.000	0.000	0.00	0	...	0	0
12214	2.484	2.484	0.00	0	...	0	0
9475	0.000	0.000	0.00	0	...	0	0

	Maharashtra	chalky	clay	loamy	peaty	sandy	silt	silty
5229	1	1	0	0	0	0	0	0
3067	1	0	1	0	0	0	0	0
6937	1	0	1	0	0	0	0	0
11148	1	0	1	0	0	0	0	0
11818	1	0	1	0	0	0	0	0
...
5269	1	0	0	0	0	0	1	0
3845	1	0	0	0	0	1	0	0
2394	1	0	0	0	0	0	1	0
12214	1	0	0	1	0	0	0	0
9475	1	1	0	0	0	0	0	0

[2177 rows x 81 columns]

8196	0.681928
1674	0.389952
7662	0.726316
2732	0.470588
12588	1.152685

6581	0.200000
9120	1.836364
9155	1.142857
9319	0.666667
1582	0.936937

Name: Yield, Length: 5078, dtype: float64

5229	0.500000
3067	1.165605
6937	0.666667
11148	0.800000
11818	0.333333

5269	0.658131
3845	1.211538
2394	0.400000
12214	0.779221
9475	0.727273

Name: Yield, Length: 2177, dtype: float64

```
In [36]: 1 import numpy as np
          2 import matplotlib.pyplot as plt
          3 import seaborn as seabornInstance
          4 from sklearn.linear_model import LinearRegression
          5 from sklearn import metrics
          6 %matplotlib inline
```

```
In [37]: 1 from sklearn.preprocessing import StandardScaler
          2 sc = StandardScaler()
          3 a_train = sc.fit_transform(a_train)
          4 a_test = sc.transform(a_test)
```

Random Forest Regressor

```
In [38]: 1 from sklearn.ensemble import RandomForestRegressor
          2 regr = RandomForestRegressor(max_depth=2, random_state=0, n_estimators=100)
          3 regr.fit(a_train, b_train)
          4 b_pred = regr.predict(a_test)
          5
          6 from sklearn.metrics import mean_squared_error as mse
          7 from sklearn.metrics import mean_absolute_error as mae
          8 from sklearn.metrics import r2_score
          9
          10 print('MSE =', mse(b_pred, b_test))
          11 print('MAE =', mae(b_pred, b_test))
          12 print('R2 Score =', r2_score(b_pred, b_test))
```

MSE = 7.67104887996405
MAE = 0.8953650873829122
R2 Score = 0.9589614680509004

Polynomial Support Vector Machine

```
In [39]: 1 from sklearn.svm import SVR
          2 regressorpoly = SVR(kernel='poly', epsilon=1.0)
          3 regressorpoly.fit(a_train, b_train)
          4 pred = regressorpoly.predict(a_test)
          5 print(regressorpoly.score(a_test, b_test))
          6 print(r2_score(b_test, b_pred))
```

0.6312485850825429
0.9598856437260073

XGBRegressor

```
In [42]: 1 from xgboost import XGBRegressor
2 from sklearn.metrics import mean_absolute_error
3 XGBModel = XGBRegressor()
4 XGBModel.fit(a_train,b_train , verbose=False)
5
6 # Get the mean absolute error on the validation data :
7 XGBpredictions = XGBModel.predict(a_test)
8 MAE = mean_absolute_error(b_test , XGBpredictions)
9 print('XGBoost validation MAE = ',MAE)
10 XGBpredictions
```

XGBoost validation MAE = 0.6670485576475581

Out[42]: array([1.0828081, 0.64165634, 0.80906236, ..., 1.2118115, 0.9264982 ,
0.68652374], dtype=float32)

```
In [43]: 1
2 print(r2_score(b_test , XGBpredictions))
```

0.9654928330252374

```
In [ ]: 1
```